

GUDIVADA VENKATA SSHA SAI DEEPAK - 21BCE9822

```
In [102]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [103]: data=pd.read_csv("Employee-Attrition.csv")
```

```
In [104]: data.head()
```

Out[104]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education-Field
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life Sciences
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life Sciences
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	Life Sciences
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life Sciences
4	27	No	Travel_Rarely	591	Research & Development	2	1	Life Sciences

5 rows × 35 columns

```
In [105]: data.tail()
```

Out[105]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Education-Field
1465	36	No	Travel_Frequently	884	Research & Development	23	2	Life Sciences
1466	39	No	Travel_Rarely	613	Research & Development	6	1	Life Sciences
1467	27	No	Travel_Rarely	155	Research & Development	4	3	Life Sciences
1468	49	No	Travel_Frequently	1023	Sales	2	3	Life Sciences
1469	34	No	Travel_Rarely	628	Research & Development	8	3	Life Sciences

5 rows × 35 columns

```
In [106]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 35 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                   1470 non-null   int64
1   Attrition                           1470 non-null   object
2   BusinessTravel                      1470 non-null   object
3   DailyRate                           1470 non-null   int64
4   Department                          1470 non-null   object
5   DistanceFromHome                   1470 non-null   int64
6   Education                           1470 non-null   int64
7   EducationField                      1470 non-null   object
8   EmployeeCount                       1470 non-null   int64
9   EmployeeNumber                     1470 non-null   int64
10  EnvironmentSatisfaction              1470 non-null   int64
11  Gender                              1470 non-null   object
12  HourlyRate                           1470 non-null   int64
13  JobInvolvement                      1470 non-null   int64
14  JobLevel                            1470 non-null   int64
15  JobRole                             1470 non-null   object
16  JobSatisfaction                     1470 non-null   int64
17  MaritalStatus                       1470 non-null   object
18  MonthlyIncome                       1470 non-null   int64
19  MonthlyRate                         1470 non-null   int64
20  NumCompaniesWorked                  1470 non-null   int64
21  Over18                              1470 non-null   object
22  OverTime                            1470 non-null   object
23  PercentSalaryHike                   1470 non-null   int64
24  PerformanceRating                   1470 non-null   int64
25  RelationshipSatisfaction             1470 non-null   int64
26  StandardHours                       1470 non-null   int64
27  StockOptionLevel                    1470 non-null   int64
28  TotalWorkingYears                   1470 non-null   int64
29  TrainingTimesLastYear               1470 non-null   int64
30  WorkLifeBalance                     1470 non-null   int64
31  YearsAtCompany                      1470 non-null   int64
32  YearsInCurrentRole                  1470 non-null   int64
33  YearsSinceLastPromotion              1470 non-null   int64
34  YearsWithCurrManager                1470 non-null   int64
dtypes: int64(26), object(9)
memory usage: 402.1+ KB
```

```
In [107]: data.describe()
```

Out[107]:

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNu
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.00
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.86
std	9.135373	403.509100	8.106864	1.024165	0.0	602.02
min	18.000000	102.000000	1.000000	1.000000	1.0	1.00
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.25
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.50
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.75
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.00

8 rows × 26 columns

Handling Null Values

```
In [108]: data.isnull().any()
```

```
Out[108]: Age                False
Attrition                  False
BusinessTravel             False
DailyRate                 False
Department                False
DistanceFromHome          False
Education                 False
EducationField             False
EmployeeCount             False
EmployeeNumber            False
EnvironmentSatisfaction   False
Gender                    False
HourlyRate                False
JobInvolvement            False
JobLevel                  False
JobRole                   False
JobSatisfaction           False
MaritalStatus             False
MonthlyIncome             False
MonthlyRate              False
NumCompaniesWorked        False
Over18                    False
OverTime                  False
PercentSalaryHike         False
PerformanceRating         False
RelationshipSatisfaction  False
StandardHours            False
StockOptionLevel         False
TotalWorkingYears         False
TrainingTimesLastYear    False
WorkLifeBalance           False
YearsAtCompany            False
YearsInCurrentRole        False
YearsSinceLastPromotion   False
YearsWithCurrManager      False
dtype: bool
```

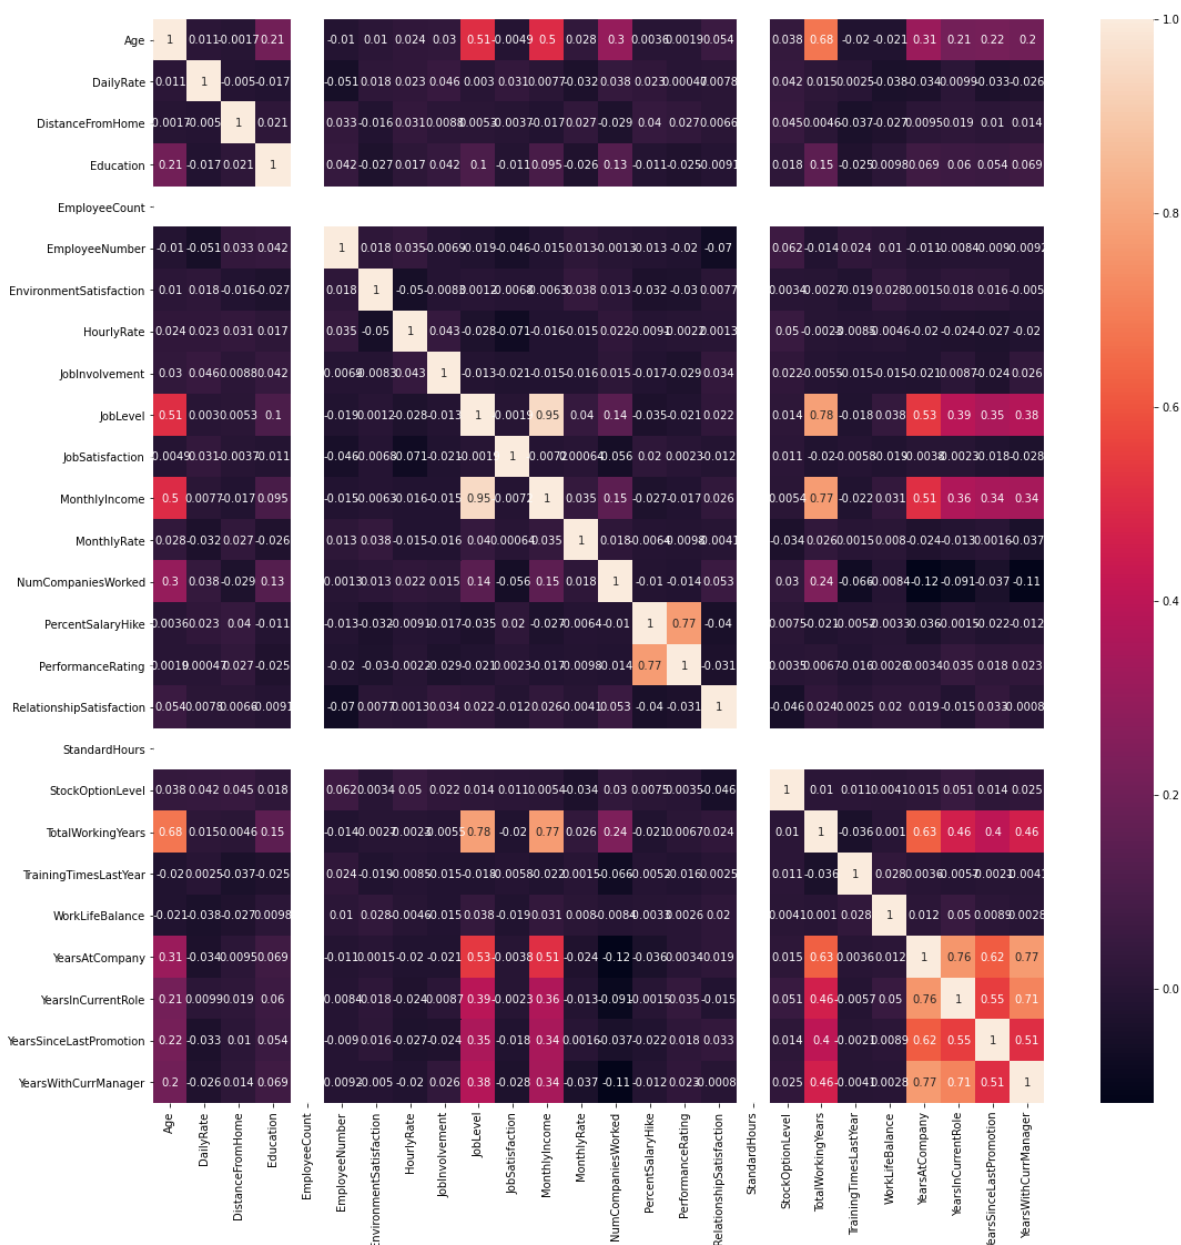
```
In [109]: data.isnull().sum()
```

```
Out[109]: Age                                0
Attrition                                   0
BusinessTravel                             0
DailyRate                                  0
Department                                 0
DistanceFromHome                           0
Education                                   0
EducationField                              0
EmployeeCount                               0
EmployeeNumber                             0
EnvironmentSatisfaction                    0
Gender                                      0
HourlyRate                                  0
JobInvolvement                             0
JobLevel                                    0
JobRole                                     0
JobSatisfaction                            0
MaritalStatus                              0
MonthlyIncome                              0
MonthlyRate                                 0
NumCompaniesWorked                         0
Over18                                      0
OverTime                                    0
PercentSalaryHike                          0
PerformanceRating                          0
RelationshipSatisfaction                    0
StandardHours                              0
StockOptionLevel                           0
TotalWorkingYears                         0
TrainingTimesLastYear                      0
WorkLifeBalance                            0
YearsAtCompany                             0
YearsInCurrentRole                         0
YearsSinceLastPromotion                    0
YearsWithCurrManager                       0
dtype: int64
```

```
In [110]: cor=data.corr()
```

```
In [111]: fig=plt.figure(figsize=(18,18))
sns.heatmap(cor,annot=True)
```

```
Out[111]: <AxesSubplot:>
```




Outliers

```
In [112]: sns.boxplot(data[ "Age" ])
```

```
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From v  
ersion 0.12, the only valid positional argument will be `data`, and pa  
ssing other arguments without an explicit keyword will result in an er  
ror or misinterpretation.  
warnings.warn(  

```

```
Out[112]: <AxesSubplot:xlabel='Age'>
```




```
In [113]: sns.boxplot(data[ "DailyRate" ])
```

```
/opt/anaconda3/lib/python3.8/site-packages/seaborn/_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From v  
ersion 0.12, the only valid positional argument will be `data`, and pa  
ssing other arguments without an explicit keyword will result in an er  
ror or misinterpretation.  
warnings.warn(  

```

```
Out[113]: <AxesSubplot:xlabel='DailyRate'>
```



```
In [114]: data.describe()
```

```
Out[114]:
```

	Age	DailyRate	DistanceFromHome	Education	EmployeeCount	EmployeeNu
count	1470.000000	1470.000000	1470.000000	1470.000000	1470.0	1470.00
mean	36.923810	802.485714	9.192517	2.912925	1.0	1024.86
std	9.135373	403.509100	8.106864	1.024165	0.0	602.02
min	18.000000	102.000000	1.000000	1.000000	1.0	1.00
25%	30.000000	465.000000	2.000000	2.000000	1.0	491.25
50%	36.000000	802.000000	7.000000	3.000000	1.0	1020.50
75%	43.000000	1157.000000	14.000000	4.000000	1.0	1555.75
max	60.000000	1499.000000	29.000000	5.000000	1.0	2068.00

8 rows × 26 columns

```
In [115]: data.head()
```


```
Out[115]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educ
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns



```
In [116]: fig, axes = plt.subplots(2,2)
sns.boxplot(data=data[ "YearsInCurrentRole" ],ax=axes[ 0,0 ])
sns.boxplot(data=data[ "YearsSinceLastPromotion" ],ax=axes[ 0,1 ])
sns.boxplot(data=data[ "YearsWithCurrManager" ],ax=axes[ 1,0 ])
sns.boxplot(data=data[ "WorkLifeBalance" ],ax=axes[ 1,1 ])
```

Out[116]: <AxesSubplot:>



```
In [117]: fig, axes = plt.subplots(2,2)
sns.boxplot(data=data[ "DistanceFromHome" ],ax=axes[ 0,0 ])
sns.boxplot(data=data[ "TotalWorkingYears" ],ax=axes[ 0,1 ])
sns.boxplot(data=data[ "HourlyRate" ],ax=axes[ 1,0 ])
sns.boxplot(data=data[ "YearsAtCompany" ],ax=axes[ 1,1 ])
```

Out[117]: <AxesSubplot:>



Handling the Outliers

```
In [118]: YearsInCurrentRole_q1 = data.YearsInCurrentRole.quantile(0.25)
YearsInCurrentRole_q3 = data.YearsInCurrentRole.quantile(0.75)
IQR_YearsInCurrentRole=YearsInCurrentRole_q3-YearsInCurrentRole_q1
upperlimit_YearsInCurrentRole=YearsInCurrentRole_q3+1.5*IQR_YearsInCurrentRole
lower_limit_YearsInCurrentRole =YearsInCurrentRole_q1-1.5*IQR_YearsInCurrentRole
median_YearsInCurrentRole=data["YearsInCurrentRole"].median()
data['YearsInCurrentRole'] = np.where(
    (data['YearsInCurrentRole'] > upperlimit_YearsInCurrentRole),
    median_YearsInCurrentRole,
    data['YearsInCurrentRole']
)
```

```
In [119]: YearsSinceLastPromotion_q1 = data.YearsSinceLastPromotion.quantile(0.25)
YearsSinceLastPromotion_q3 = data.YearsSinceLastPromotion.quantile(0.75)
IQR_YearsSinceLastPromotion=YearsSinceLastPromotion_q3-YearsSinceLastPromotion_q1
upperlimit_YearsSinceLastPromotion=YearsSinceLastPromotion_q3+1.5*IQR_YearsSinceLastPromotion
lower_limit_YearsSinceLastPromotion =YearsSinceLastPromotion_q1-1.5*IQR_YearsSinceLastPromotion
median_YearsSinceLastPromotion=data["YearsSinceLastPromotion"].median()
data['YearsSinceLastPromotion'] = np.where(
    (data['YearsSinceLastPromotion'] > upperlimit_YearsSinceLastPromotion),
    median_YearsSinceLastPromotion,
    data['YearsSinceLastPromotion']
)
```


```
In [120]: YearsWithCurrManager_q1 = data.YearsWithCurrManager.quantile(0.25)
YearsWithCurrManager_q3 = data.YearsWithCurrManager.quantile(0.75)
IQR_YearsWithCurrManager=YearsWithCurrManager_q3-YearsWithCurrManager_q1
upperlimit_YearsWithCurrManager=YearsWithCurrManager_q3+1.5*IQR_YearsWithCurrManager
lower_limit_YearsWithCurrManager =YearsWithCurrManager_q1-1.5*IQR_YearsWithCurrManager
median_YearsWithCurrManager=data["YearsWithCurrManager"].median()
data['YearsWithCurrManager'] = np.where(
    (data['YearsWithCurrManager'] > upperlimit_YearsWithCurrManager),
    median_YearsWithCurrManager,
    data['YearsWithCurrManager']
)
```

```
In [121]: TotalWorkingYears_q1 = data.TotalWorkingYears.quantile(0.25)
TotalWorkingYears_q3 = data.TotalWorkingYears.quantile(0.75)
IQR_TotalWorkingYears=TotalWorkingYears_q3-TotalWorkingYears_q1
upperlimit_TotalWorkingYears=TotalWorkingYears_q3+1.5*IQR_TotalWorkingYears
lower_limit_TotalWorkingYears=TotalWorkingYears_q1-1.5*IQR_TotalWorkingYears
median_TotalWorkingYears=data["TotalWorkingYears"].median()
data['TotalWorkingYears'] = np.where(
    (data['TotalWorkingYears'] > upperlimit_TotalWorkingYears),
    median_TotalWorkingYears,
    data['TotalWorkingYears']
)
```

```
In [122]: YearsAtCompany_q1 = data.YearsAtCompany.quantile(0.25)
YearsAtCompany_q3 = data.YearsAtCompany.quantile(0.75)
IQR_YearsAtCompany=YearsAtCompany_q3-YearsAtCompany_q1
upperlimit_YearsAtCompany=YearsAtCompany_q3+1.5*IQR_YearsAtCompany
lower_limit_YearsAtCompany=YearsAtCompany_q1-1.5*IQR_YearsAtCompany
median_YearsAtCompany=data["YearsAtCompany"].median()
data['YearsAtCompany'] = np.where(
    (data['YearsAtCompany'] > upperlimit_YearsAtCompany),
    median_YearsAtCompany,
    data['YearsAtCompany']
)
```

```
In [123]: fig, axes = plt.subplots(2,2)
sns.boxplot(data=data["YearsWithCurrManager"],ax=axes[0,0])
sns.boxplot(data=data["TotalWorkingYears"],ax=axes[0,1])
sns.boxplot(data=data["YearsSinceLastPromotion"],ax=axes[1,0])
sns.boxplot(data=data["YearsAtCompany"],ax=axes[1,1])
```

Out[123]: <AxesSubplot:>



```
In [124]: data.head()
```

Out[124]:

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educ
0	41	Yes	Travel_Rarely	1102	Sales	1	2	Life
1	49	No	Travel_Frequently	279	Research & Development	8	1	Life
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	Life
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 35 columns

```
In [125]: data.drop("EducationField",axis=1,inplace=True)
```

```
In [126]: data.head()
```

```
Out[126]:
```

	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Empl
0	41	Yes	Travel_Rarely	1102	Sales	1	2	
1	49	No	Travel_Frequently	279	Research & Development	8	1	
2	37	Yes	Travel_Rarely	1373	Research & Development	2	2	
3	33	No	Travel_Frequently	1392	Research & Development	3	4	
4	27	No	Travel_Rarely	591	Research & Development	2	1	

5 rows × 34 columns

```
In [127]: data["BusinessTravel"].unique()
```

```
Out[127]: array(['Travel_Rarely', 'Travel_Frequently', 'Non-Travel'], dtype=object)
```

Splitting the data

```
In [128]: y=data["Attrition"]
```

```
In [129]: y.head()
```

```
Out[129]: 0    Yes
          1    No
          2    Yes
          3    No
          4    No
          Name: Attrition, dtype: object
```

```
In [130]: data.drop("Attrition",axis=1,inplace=True)
```

```
In [131]: data.head()
```

```
Out[131]:
```

	Age	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	EmployeeCount
0	41	Travel_Rarely	1102	Sales	1	2	1
1	49	Travel_Frequently	279	Research & Development	8	1	1
2	37	Travel_Rarely	1373	Research & Development	2	2	1
3	33	Travel_Frequently	1392	Research & Development	3	4	1
4	27	Travel_Rarely	591	Research & Development	2	1	1

5 rows × 33 columns

Encoding

```
In [132]: from sklearn.preprocessing import LabelEncoder
```

```
In [133]: le=LabelEncoder()
```

```
In [134]: data["BusinessTravel"]=le.fit_transform(data["BusinessTravel"])
```

```
In [135]: data["Department"]=le.fit_transform(data["Department"])
```

```
In [136]: data["Gender"]=le.fit_transform(data["Gender"])
```

```
In [137]: y=le.fit_transform(y)
```

```
In [138]: y
```

```
Out[138]: array([1, 0, 1, ..., 0, 0, 0])
```

```
In [139]: data["JobRole"]=le.fit_transform(data["JobRole"])
```

```
In [140]: data["Over18"]=le.fit_transform(data["Over18"])
```

```
In [141]: data["MaritalStatus"]=le.fit_transform(data["MaritalStatus"])
```

```
In [142]: data["OverTime"]=le.fit_transform(data["OverTime"])
```

```
In [143]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1470 entries, 0 to 1469
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                    1470 non-null   int64
1   BusinessTravel                       1470 non-null   int64
2   DailyRate                           1470 non-null   int64
3   Department                          1470 non-null   int64
4   DistanceFromHome                    1470 non-null   int64
5   Education                           1470 non-null   int64
6   EmployeeCount                       1470 non-null   int64
7   EmployeeNumber                      1470 non-null   int64
8   EnvironmentSatisfaction              1470 non-null   int64
9   Gender                              1470 non-null   int64
10  HourlyRate                          1470 non-null   int64
11  JobInvolvement                      1470 non-null   int64
12  JobLevel                            1470 non-null   int64
13  JobRole                             1470 non-null   int64
14  JobSatisfaction                     1470 non-null   int64
15  MaritalStatus                      1470 non-null   int64
16  MonthlyIncome                      1470 non-null   int64
17  MonthlyRate                         1470 non-null   int64
18  NumCompaniesWorked                  1470 non-null   int64
19  Over18                             1470 non-null   int64
20  OverTime                           1470 non-null   int64
21  PercentSalaryHike                   1470 non-null   int64
22  PerformanceRating                   1470 non-null   int64
23  RelationshipSatisfaction             1470 non-null   int64
24  StandardHours                      1470 non-null   int64
25  StockOptionLevel                    1470 non-null   int64
26  TotalWorkingYears                   1470 non-null   float64
27  TrainingTimesLastYear               1470 non-null   int64
28  WorkLifeBalance                     1470 non-null   int64
29  YearsAtCompany                      1470 non-null   float64
30  YearsInCurrentRole                  1470 non-null   float64
31  YearsSinceLastPromotion              1470 non-null   float64
32  YearsWithCurrManager                 1470 non-null   float64
dtypes: float64(5), int64(28)
memory usage: 379.1 KB
```

Train Test Split

```
In [144]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(data,y,test_size=0.3,rand
```

```
In [145]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[145]: ((1029, 33), (441, 33), (1029,), (441,))
```

Featuring Scaling

```
In [146]: from sklearn.preprocessing import StandardScaler
```

```
In [147]: sc=StandardScaler()
```

```
In [148]: x_train=sc.fit_transform(x_train)
```

```
In [149]: x_test=sc.fit_transform(x_test)
```

Building the model

Multi Linear Regression

```
In [150]: from sklearn.linear_model import LinearRegression
```

```
In [151]: lr = LinearRegression()
```

```
In [152]: lr.fit(x_train,y_train)
```

```
Out[152]: LinearRegression()
```

```
In [153]: lr.coef_ #slope(m)
```

```
Out[153]: array([-3.54940447e-02,  7.88352347e-05, -1.70825038e-02,  3.46389690e-02,
                2.44612841e-02,  3.65668214e-03, -2.50667542e-16, -9.46820520e-03,
                -4.11203734e-02,  1.06338881e-02, -2.97662154e-03, -3.84864283e-02,
                -1.52927977e-02, -1.57839139e-02, -3.67252862e-02,  3.35765928e-02,
                -5.90043558e-03,  5.81099165e-03,  3.78471890e-02, -6.93889390e-18,
                9.55263279e-02, -2.55800078e-02,  2.01844797e-02, -2.64773510e-02,
                2.60208521e-18, -1.79286106e-02, -3.30529386e-02, -1.09247807e-02,
                -3.10631611e-02, -2.47887717e-02, -1.10177742e-02,  2.11897289e-02,
                -6.60823991e-03])
```

```
In [154]: lr.intercept_  #(c)
```

```
Out[154]: 0.16229348882410102
```

```
In [155]: y_pred = lr.predict(x_test)
```

```
In [156]: y_pred
```

```
-0.1,      6.66728668e-02,  4.49620331e-02,  3.30502696e-01,  9.74393000e
-02,      5.51447175e-01,  1.52212203e-01,  3.58819339e-01,  3.66371593e
-01,      2.47091987e-01,  5.86970935e-02,  1.28678988e-01,  2.80584025e
-01,      7.21059443e-02, -8.07006907e-02,  3.39791632e-01,  8.25270203e
-02,      2.20338157e-01,  2.47703594e-01,  4.97067397e-01,  1.36010592e
-01,      2.88153807e-01,  4.61306498e-02,  4.52544344e-01, -8.24037634e
-02,      2.26796295e-01,  1.42129836e-02,  1.62111340e-01,  2.32246950e
-01,      9.12503556e-02,  1.18866795e-01,  2.12735292e-01, -2.69559828e
-02,      4.53611463e-02,  1.09618223e-01,  2.64436901e-02,  2.32180310e
-01,      1.63285101e-01,  2.42669261e-01,  5.44757533e-01,  1.25881866e
~ -
```



```
In [157]: y_test
```

```
Out[157]: array([0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
          0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
          1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1,
          1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
1,
          0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
          0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
          1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
0,
          0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
          0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
          1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
1,
          0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0])
```

Logistic Regression

```
In [158]: from sklearn.linear_model import LogisticRegression
```

```
In [159]: lg=LogisticRegression()
```

```
In [160]: lg.fit(x_train,y_train)
```

```
Out[160]: LogisticRegression()
```

```
In [161]: y_pred_lg=lg.predict(x_test)
```

```
In [162]: y_pred
```

```
Out[162]: array([ 1.30302477e-01,  2.17626230e-01,  3.46282415e-01,  5.41382549e
-03,
         4.99292896e-01,  1.01628868e-01,  3.44742777e-01,  1.23994945e
-01,
        -1.60694945e-01,  4.02435622e-01,  1.44159172e-01,  2.67416840e
-01,
        -4.62559536e-02,  5.58671849e-01,  2.81858700e-01,  1.53537792e
-02,
         1.78573363e-01,  2.77532834e-01,  9.37121052e-02,  2.17571624e
-01,
         2.65936178e-01,  1.41499184e-02,  8.36251186e-02,  9.58849826e
-02,
         5.09869963e-01,  2.94764240e-01,  7.85819529e-02,  1.26647773e
-01,
         5.05518902e-01,  8.48456917e-02, -7.97229275e-02,  2.15516993e
-02,
         1.08079105e-01,  3.65998400e-01,  1.24517362e-01,  5.13682786e
-02,
         1.06749689e-01,  6.07640778e-02,  6.66425313e-02,  4.81312859e
-02])
```

```
In [163]: y_test
```

```
Out[163]: array([0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
        0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
        1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1,
        1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
1,
        0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
        0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
        1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0,
        0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0,
        0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
        0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
        0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
        0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
        1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
        0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
1,
        0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
        0])
```

```
In [164]: score = lg.score(x_test, y_test)
print(score)
```

```
0.8820861678004536
```

Confusion Matrix

```
In [165]: from sklearn import metrics
cm = metrics.confusion_matrix(y_test, y_pred_lg)
print(cm)
```

```
[[366   5]
 [ 47  23]]
```

Ridge and Lasso

```
In [166]: from sklearn.linear_model import Ridge
          from sklearn.model_selection import GridSearchCV
```

```
In [167]: rg=Ridge()
```

```
In [168]: parameters={"alpha":[1,2,3,5,10,20,30,40,60,70,80,90]}
          ridgecv=GridSearchCV(rg,parameters,scoring="neg_mean_squared_error",cv=5)
          ridgecv.fit(x_train,y_train)
```

```
Out[168]: GridSearchCV(cv=5, estimator=Ridge(),
                      param_grid={'alpha': [1, 2, 3, 5, 10, 20, 30, 40, 60, 70,
                      80, 90]},
                      scoring='neg_mean_squared_error')
```

```
In [169]: print(ridgecv.best_params_)

{'alpha': 90}
```

```
In [170]: print(ridgecv.best_score_)

-0.11390621139234183
```

```
In [171]: y_pred_rg=ridgecv.predict(x_test)
```

```
In [172]: y_pred_rg
```

```
Out[172]: array([[ 1.34413485e-01,  2.22561818e-01,  3.41692977e-01,  3.88209867e
-03,
                  4.84617338e-01,  1.16361483e-01,  3.30449743e-01,  1.27358807e
-01,
                 -1.34442619e-01,  3.77692888e-01,  1.33001445e-01,  2.69898751e
-01,
                 -2.54707392e-02,  5.25771894e-01,  2.67543514e-01,  2.78725024e
-02,
                  1.82233111e-01,  2.78896415e-01,  9.12689699e-02,  2.11494641e
-01,
                  2.70103341e-01,  8.44922044e-03,  8.74746722e-02,  1.05348798e
-01,
                  4.87749940e-01,  2.83080512e-01,  8.80556209e-02,  1.23817268e
-01,
                  4.82185624e-01,  9.34824523e-02, -7.16448509e-02,  4.07003104e
-02,
                  1.08437994e-01,  3.42151399e-01,  1.22270929e-01,  6.85889862e
-02,
                  1.06690533e-01,  7.08689637e-02,  7.51570276e-02,  6.05829413e
-02])
```

```
In [173]: y_test
```

```
Out[173]: array([[0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
1,
0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0])
```

```
In [174]: from sklearn import metrics
print(metrics.r2_score(y_test,y_pred_rg))
print(metrics.r2_score(y_train,ridgecv.predict(x_train)))
```

```
0.21073458438815873
```

```
0.2061567210285108
```

Lasso

```
In [175]: from sklearn.linear_model import Lasso
from sklearn.model_selection import GridSearchCV
```

```
In [176]: la=Ridge()
```

```
In [177]: parametres={"alpha":[1,2,3,5,10,20,30,40,60,70,80,90]}
ridgecv=GridSearchCV(la,parametres,scoring="neg_mean_squared_error",cv=5)
ridgecv.fit(x_train,y_train)
```

```
Out[177]: GridSearchCV(cv=5, estimator=Ridge(),
                        param_grid={'alpha': [1, 2, 3, 5, 10, 20, 30, 40, 60, 70,
80, 90]},
                        scoring='neg_mean_squared_error')
```

```
In [178]: print(ridgecv.best_params_)

{'alpha': 90}
```

```
In [179]: print(ridgecv.best_score_)

-0.11390621139234183
```

```
In [180]: y_pred_la=ridgecv.predict(x_test)
```

```
In [181]: y_pred_la
          7.50317322e-02,  1.67646673e-01,  1.16585544e-01,  1.07157808e
-01,
        -1.84689359e-02,  1.86217544e-01,  1.16586463e-01,  4.67201201e
-02,
          1.11060472e-01,  2.27053971e-01, -7.00247692e-02, -5.81070776e
-02,
          2.03141688e-01,  4.69029664e-02,  1.31525768e-01,  5.66738022e
-01,
          2.41883060e-02, -3.41250985e-02, -1.13904557e-01,  2.18572744e
-01,
          2.60568042e-01,  1.65533667e-01, -5.94078459e-05,  2.60009384e
-01,
          4.20709666e-01,  3.71031267e-01,  1.70250288e-01,  4.03052216e
-01,
          4.67312765e-01,  1.98845366e-01,  1.55005619e-01,  3.41505080e
-01,
          2.20024496e-01,  1.40989758e-01,  1.97796963e-01,  2.57841889e
-01,
          2.99122317e-01,  9.24907038e-03,  1.39162817e-01, -1.13916709e
-01,
```

```
In [182]: from sklearn import metrics
print(metrics.r2_score(y_test,y_pred_la))
print(metrics.r2_score(y_train,ridgecv.predict(x_train)))

0.21073458438815873
0.2061567210285108
```

Decision Tree

```
In [183]: from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
```

```
In [184]: dtc.fit(x_train,y_train)
```

```
Out[184]: DecisionTreeClassifier()
```

```
In [185]: pred=dtc.predict(x_test)
```

```
In [186]: pred
```

```
Out[186]: array([0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0,
0,
          0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0,
          0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
          0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
          0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
1,
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1,
0,
          0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1,
1,
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0,
0,
          0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0,
0,
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
1,
          0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
          0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1,
0,
          0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
0,
          0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
          0])
```

In [187]: `y_test`

Out[187]: `array([0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0,
1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
0,
0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
1,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1,
0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0,
1,
0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
0])`

In [188]: `#Accuracy score
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

In [189]: `accuracy_score(y_test, pred)`

Out[189]: `0.7664399092970522`

In [190]: `confusion_matrix(y_test, pred)`

Out[190]: `array([[320, 51],
[52, 18]])`


```
In [191]: pd.crosstab(y_test,pred)
```

```
Out[191]:
```

col_0	0	1
row_0		
0	320	51
1	52	18

```
In [192]: print(classification_report(y_test,pred))
```

	precision	recall	f1-score	support
0	0.86	0.86	0.86	371
1	0.26	0.26	0.26	70
accuracy			0.77	441
macro avg	0.56	0.56	0.56	441
weighted avg	0.77	0.77	0.77	441

```
In [193]: probability=dtc.predict_proba(x_test)[: ,1]
```

```
In [194]: # roc_curve
fpr,tpr,threshholds = roc_curve(y_test,probability)
```

```
In [195]: plt.plot(fpr,tpr)
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('ROC CURVE')
plt.show()
```

